L&P Platform Scale Options

1 · THE GOALS

Five business results, and what each one requires the platform to do

PRIMARY
Stronger proposals

- Accurate, complete, compliant proposals for each RFP
- Know what was strong, what failed and what to change
- Staffing and price plans leadership can approve

Document intake · Compliance matrix · Pricing engine · Win/loss history · Price-to-win

GROWTH
Find and convert business

- Which buyers to pursue, when, and with what advantage
- A qualified next-contact list with “why now”
- Which sources and messages win customers

Market analysis · Recompete calendar · Lead scoring · Pwin · Outreach tracking

PROTECTION
Protect revenue and clients

- Authorization gaps and renewals flagged in time
- Contract economics and billing exceptions explained
- Client issues owned before renewal

Contract records · QuickBooks + Belfry · Alerts · Client reviews

READINESS
Strengthen capability

- Verified past performance ready to reuse
- Credential and eligibility gaps visible
- Partner, expansion and training plans

Evidence library · Credential tracking · Expansion tracking

SHARED
Reusable intelligence

- Every bid adds knowledge for the next
- ChatGPT and Claude work from the same records
- Less repeated research and re-entry

One source of truth · AI agent access · Automation

ALSO REQUESTED — Pwin (probability of win) · Price-to-win · Market analysis · Lead scoring · AI automation for speed and accuracy

1 · Collect	2 · Verify	3 · Analyze	4 · Decide	5 · Produce	6 · Learn
RFPs, our bids, competitor awards, market data	AI extracts, staff confirm, every fact sourced	Comparisons, Pwin, price-to-win, leads	Bid or not, staffing, price	Proposal, forms, price sheets	Outcome and feedback feed the next bid

Source: the platform Reference sheet (Overview and Features tabs) and leadership requests.

Eleven layers with the database at the center; each option builds on the one before

Sources — RFPs, past bids, public awards, QuickBooks, Belfry → **Integrations + processing** — pull, read and extract (Cloud Run) → **Database + documents** — PostgreSQL + file storage: the system of record → **Logic, AI + analytics** — pricing, scoring, BigQuery models → **Interface + AI agents** — staff, ChatGPT and Claude

Use · Intelligence · Data · Operations · Included · Partial · + Added by this option

	1 · Intelligence platform Smallest real starting point	2 · Enterprise data & AI Pwin, price-to-win, leads, market	3 · Full operating platform Everything in the Reference sheet
Interface Next.js web app · Vercel	Bid, pricing and contract workspace	+ Market and pipeline dashboards	+ Sales, client and executive views
AI agent access MCP server · sign-in	Per-person permissions; ChatGPT and Claude	Same as option 1	+ Autonomous scheduled agents
Business logic App services · workflows	Staffing and pricing engine, approvals	+ Pwin scorecard, price-to-win, lead scoring	+ Outreach sequences, records-request follow-up
AI services Gemini · Vertex AI	Extraction, drafting, meaning-based search	+ Model insights on scores	+ Agent automation at volume
Data warehouse BigQuery · Dataform	Not yet: reports run on the database	+ Multi-year history, Pwin and price models	+ Marketing, client and expansion analytics
Database PostgreSQL · Supabase	**Bids, requirements, prices, staffing, contracts, competitors, outcomes**	+ Change history, staging copy, point-in-time recovery	+ Leads, outreach, client reviews, credentials
Document storage Supabase Storage	Original RFPs, proposals, contracts, versioned	+ Public awards and bid tabulations	Same as option 2
Search index Full-text + pgvector	Keyword and meaning search over verified text	Same as option 1	Same as option 1
Integrations Scheduled Cloud Run jobs	Procurement sites, Drive, QuickBooks, Belfry	+ USAspending, SAM.gov, wage data, bid tabs	+ Apollo / ZoomInfo, Gmail, review tools
Processing Python · Cloud Run	OCR pipeline, portal crawlers, document output	+ More portals, continuous crawling	+ 24/7 crawling and enrichment
Operations CI/CD · Sentry · monitoring	Backups, monitoring, alerts	+ Staging and production	+ Agent cost and safety limits

Why the database comes first: every other layer reads from or writes to it. AI agents can't work accurately across a spreadsheet of 87 tabs (about 2.35 million tokens, larger than any AI's reading window). In a database, an agent asks a precise question and gets back only the matching records.